



# Operating Systems

Szu-Chi Chung

Department of Applied Mathematics, National Sun Yat-sen University

# What we are going to study in this semester


---

- ▶ Introduction
- ▶ Data representation and operation
  - ▶ Number Systems
  - ▶ Data Storage
  - ▶ Operations on Data
- ▶ Computer hardware
  - ▶ Computer organization
  - ▶ Computer network and internet
- ▶ Computer software
  - ▶ Programming Languages (Python)
  - ▶ Operating system
  - ▶ Algorithms
- ▶ Data organization and abstraction
  - ▶ File Structure
  - ▶ Data Structure
  - ▶ Abstract Data Type
- ▶ Advance topics
  - ▶ Security
  - ▶ Artificial intelligence (Recorded lecture)
  - ▶ Data compression (Recorded lecture)
- ▶ Not covered
  - ▶ Software engineering
  - ▶ Databases
  - ▶ Theory of computation
  - ▶ Social media and social issues

# Introduction

---

- ▶ Computer software is divided into two broad categories: the operating system and application programs
  - ▶ Application programs use computer hardware to solve users' problems
  - ▶ The operating system, on the other hand, controls access to hardware by users



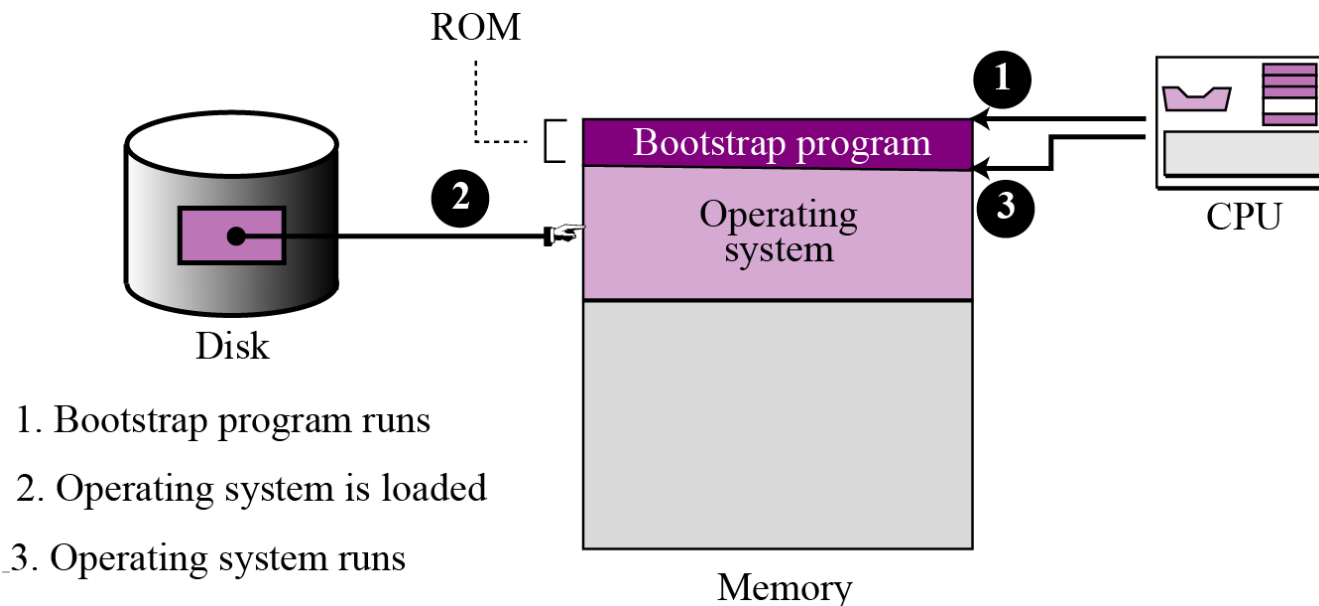
# Operating system

---

1. An operating system is an interface between the hardware of a computer and the user (programs or humans)
2. An operating system is a set of programs that facilitates the execution of other programs
3. An operating system acts as a general manager supervising the activity of each component in the computer system
  - ▶ The operating system checks that hardware and software resources are used efficiently, and when there is a conflict in using a resource, the operating system mediates to solve it
  - ▶ For example, it is responsible for loading other programs into memory for execution

# Bootstrap process

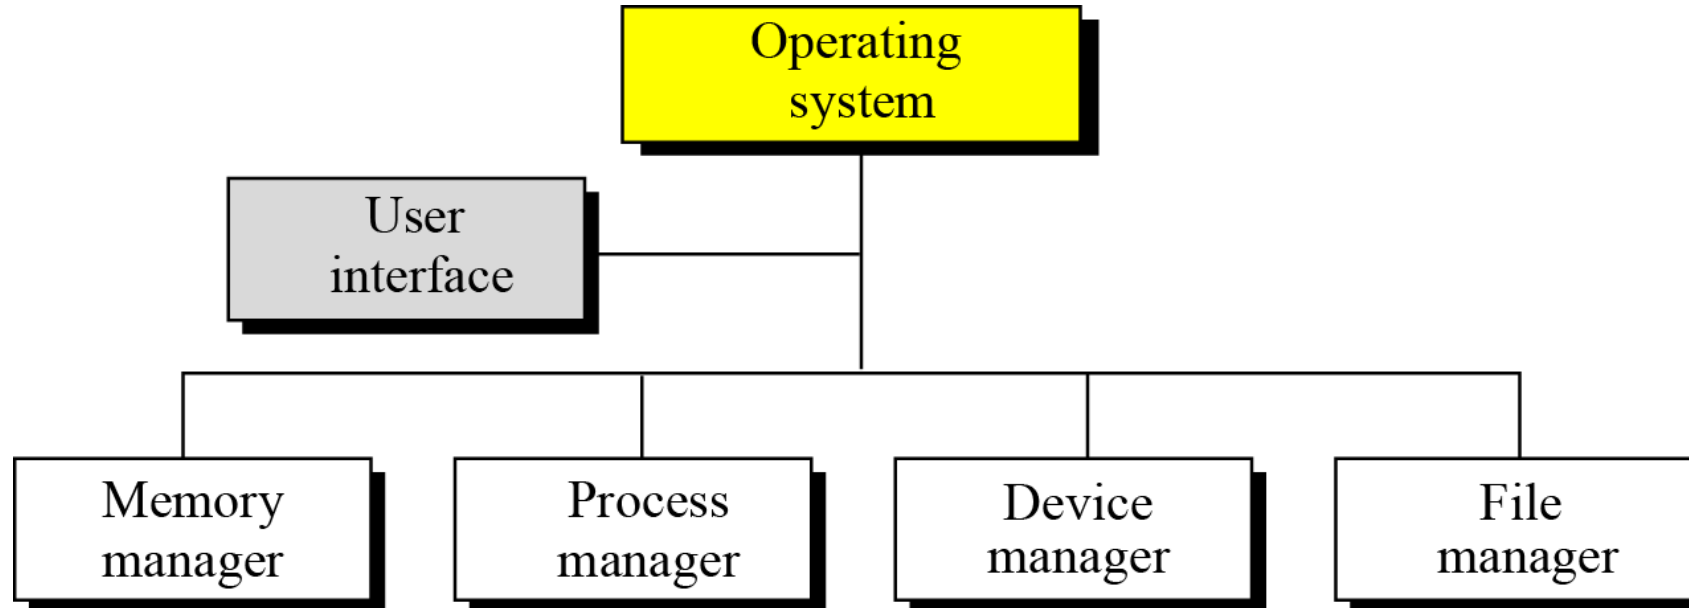
- ▶ However, the operating system itself is a program that needs to be loaded into the memory and run!
- ▶ A very small section of memory is made of ROM and holds a small program called the *bootstrap program* which is pointed by the program counter when the computer is on
- ▶ This program is only responsible for loading the operating system into RAM and the program counter in the CPU is set to the first instruction of the operating system



# Components

---


- ▶ An operating system needs to manage different resources in a computer system
  - ▶ A modern operating system has at least four duties *memory manager*, *process manager*, *device manager*, and *file manager* which are the *kernel*
  - ▶ Operating system also has a component called a user interface or a *shell* which is responsible for communication outside the operating system



# User interface

---


- ▶ User interface a program that accepts requests from users (processes) and interprets them for the rest of the operating system
  - ▶ Shells communicated with users through textual messages using a keyboard and monitor screen
  - ▶ A GUI (graphical user interface) allows users to issue commands by using a mouse or other input device to manipulate the icons
  - ▶ Touch screens allow users to manipulate icons directly with their fingers



# Memory manager

---

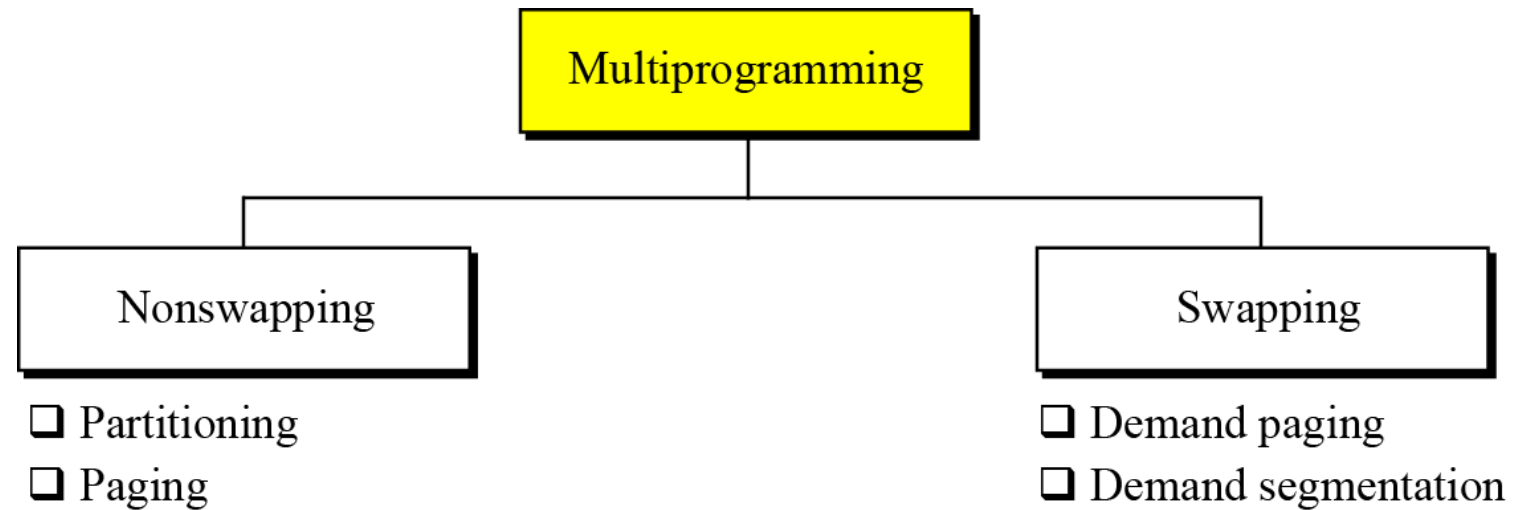
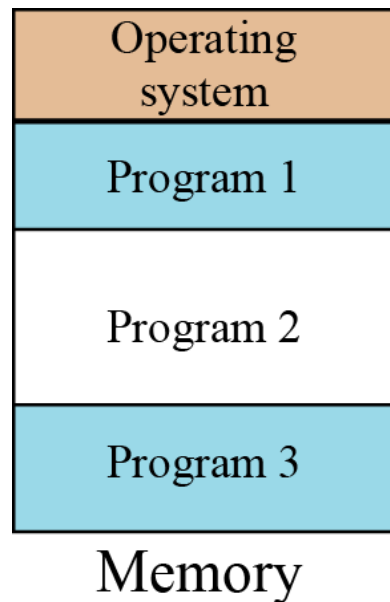
- ▶ Memory allocation must be managed to prevent applications from running out of memory
  - ▶ It can be divided into two broad categories of memory management: *monoprogramming* and *multiprogramming*
  - ▶ In monoprogramming, most of the memory capacity is dedicated to a single program
    - ▶ In this configuration, the whole program is in memory for execution
    - ▶ The job of the memory manager is to load the program into memory, run it, and replace it with the next program
    - ▶ This is a very inefficient use of memory and CPU time since we need to wait for the I/O device






# Memory manager

- ▶ In multiprogramming, more than one program is in memory at the same time
  - ▶ They are executed concurrently, with the CPU switching rapidly between the programs
  - ▶ Multiprogramming can be further divided into several categories
    - ▶ Nonswapping means that the program remains in memory for the duration of execution




# Memory manager - Partitioning

- ▶ In *partitioning*, memory is divided into variable-length partitions
  - ▶ Each section or partition holds one program. The CPU switches between programs
    - ▶ It starts with one program, executing some instructions until it either encounters an input/output operation or the time allocated for that program has expired
    - ▶ The CPU then saves the address of the memory location where the last instruction was executed and moves to the next program



a. CPU starts executing program 1



b. CPU starts executing program 2

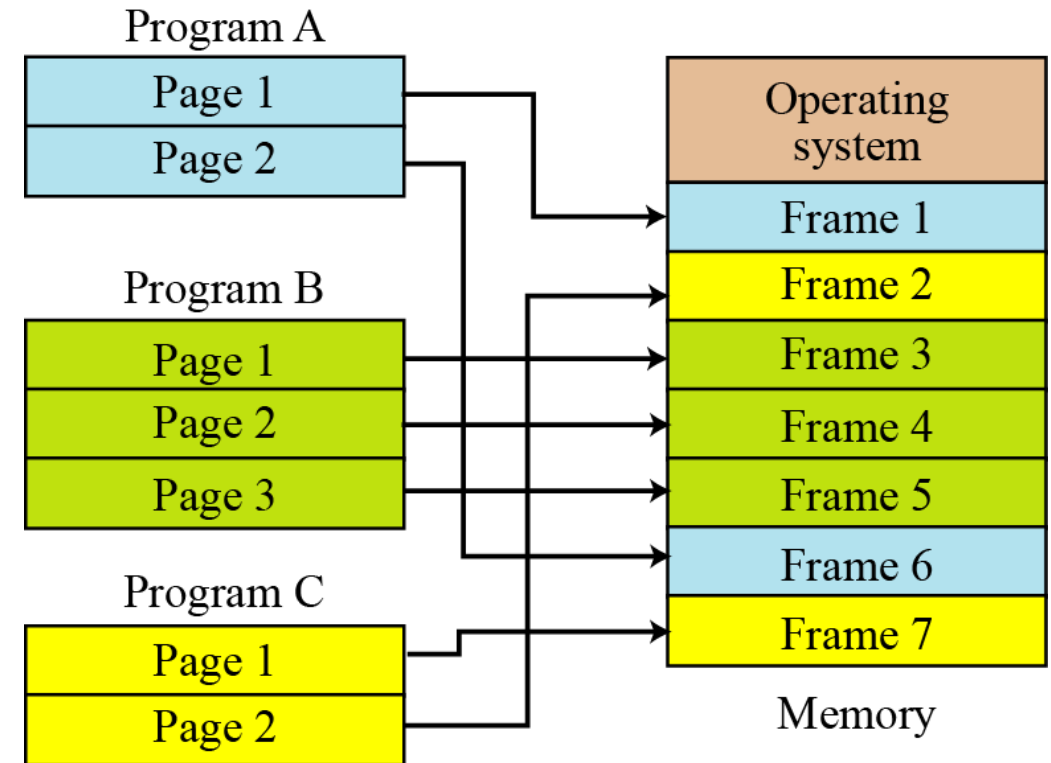
# Memory manager - Partitioning

---

- ▶ Each program is entirely in memory and occupies contiguous locations.  
However, there are some issues
  - ▶ The size of the partitions has to be determined beforehand by the memory manager
  - ▶ Even if partitioning is perfect when the computer is started, there may be some holes after completed programs are replaced by new ones
  - ▶ When there are many holes, the memory manager can compact the partitions to remove the holes and create new partitions, but this creates extra overhead on the system!

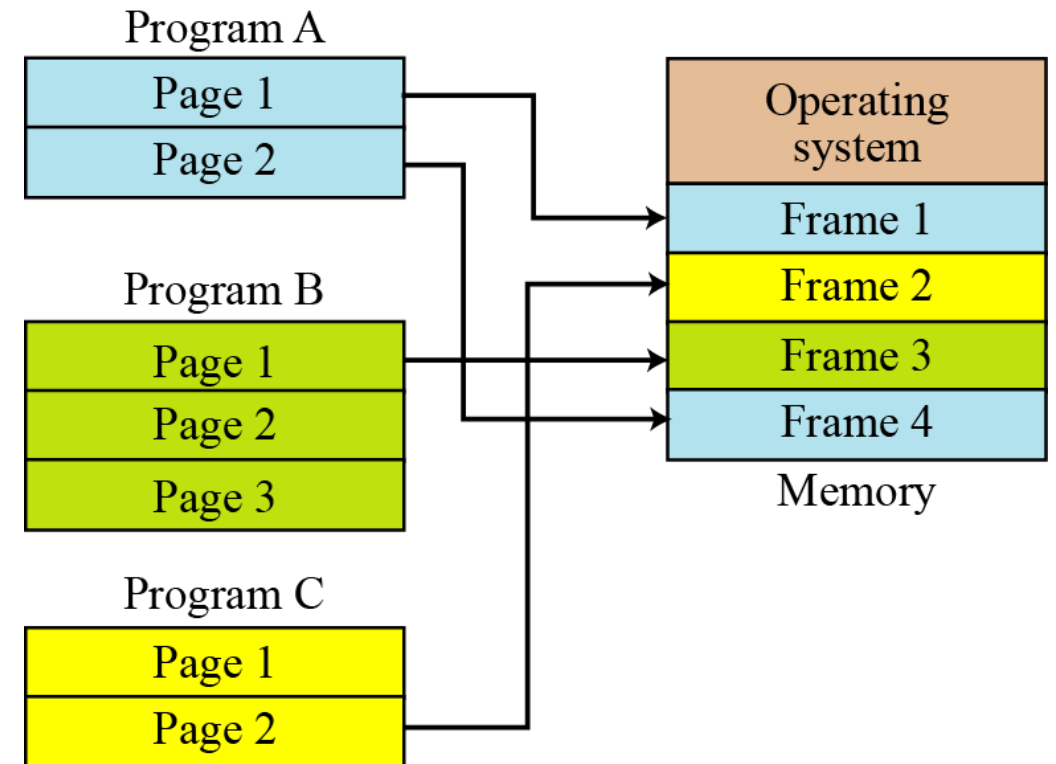
# Memory manager - Paging

- ▶ In *paging*, memory is divided into equally sized sections called frames. Programs are also divided, into equally sized sections called pages
  - ▶ The size of a page and a frame is usually the same and equal to the size of the block used by the system to retrieve information from a storage device
  - ▶ The program does not have to be contiguous in memory
  - ▶ Paging improves efficiency to some extent, but larger program still needs to be in memory before being executed



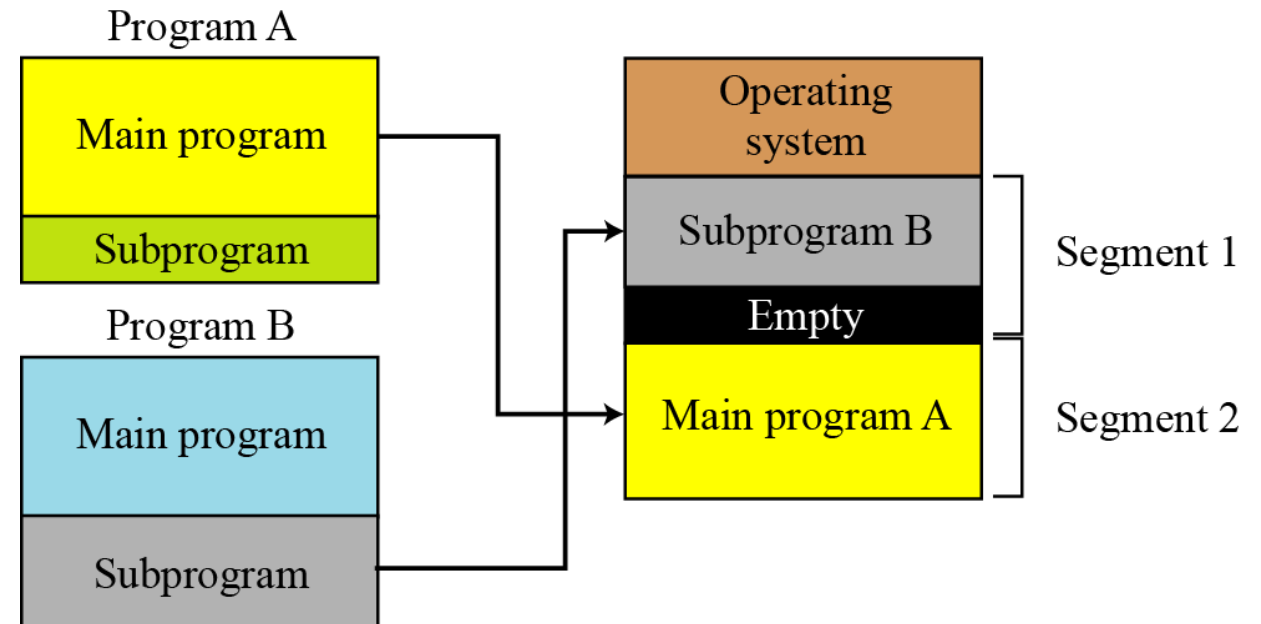
# Memory manager - Demand paging

- ▶ Paging still requires that the entire program be in memory for execution
  - ▶ In *demand paging* the program is divided into pages, but the pages can be loaded into memory one by one, executed, and replaced by another page
  - ▶ The program can be swapped between memory and disk!




# Memory manager - Demand segmentation

- ▶ In *demand segmentation*, the program is divided into segments that match the programmer's view
- ▶ In demand segmentation, the program is divided into modules that match the programmer's view
- ▶ Since segments in memory are still of equal size, part of a segment may remain empty



# Memory manager - Virtual memory


- ▶ Demand paging and demand segmentation mean that, when a program is being executed, part of the program is in memory and part is on disk!
  - ▶ This large “fictional” memory space created by paging is called *virtual memory*



# Process manager

---

- ▶ Modern operating systems use three terms that refer to a set of instructions
  - ▶ Program: A program is a nonactive set of instructions stored on a disk. It may or may not become a job
  - ▶ Job: A program becomes a job from the moment it is selected for execution until it has finished running and becomes a program again
    - ▶ During this time a job may or may not be executed and it may be on disk or in memory
  - ▶ Process: A process is a job that is being run in memory
    - ▶ As long as the job is in memory, it is a process

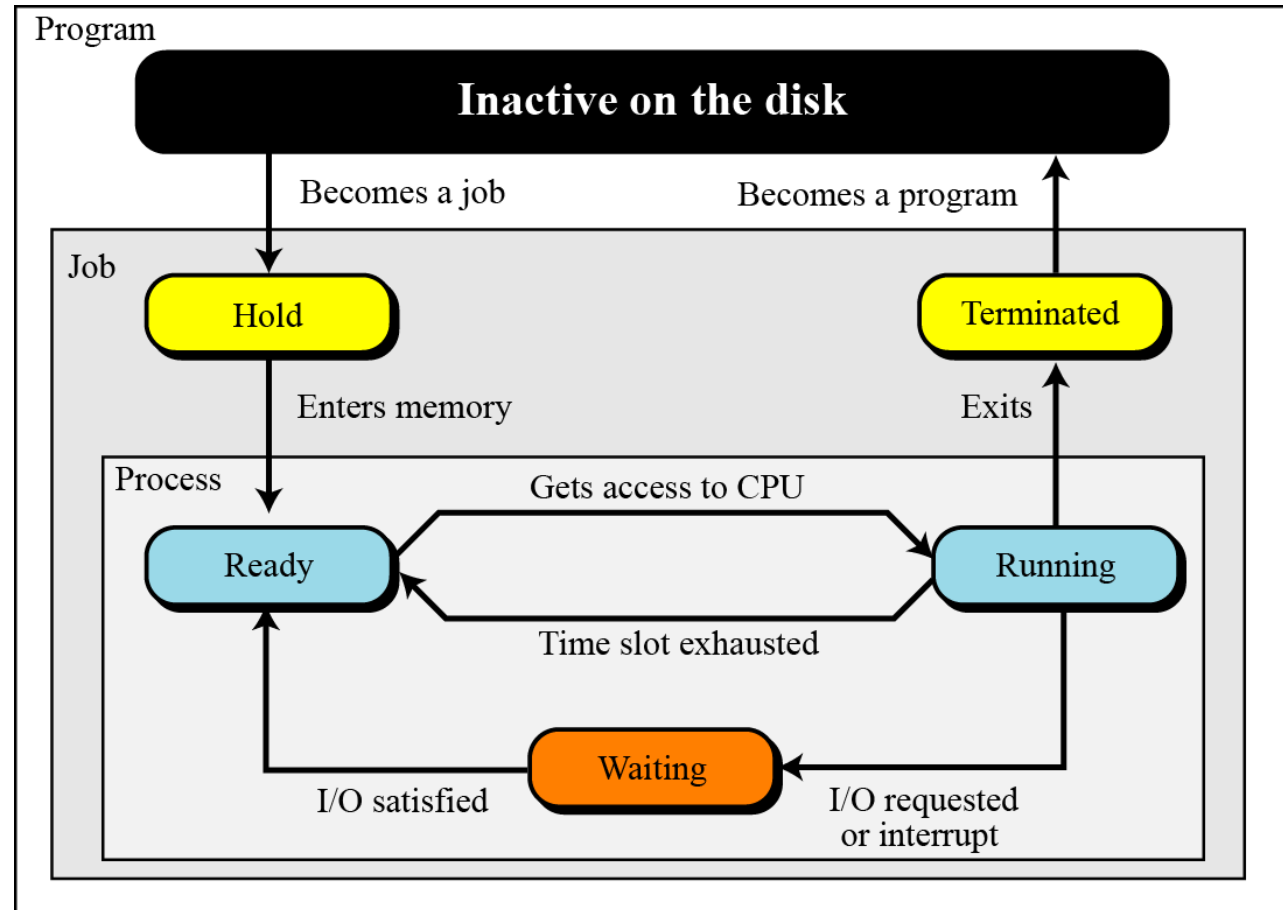




# Process manager - State diagrams

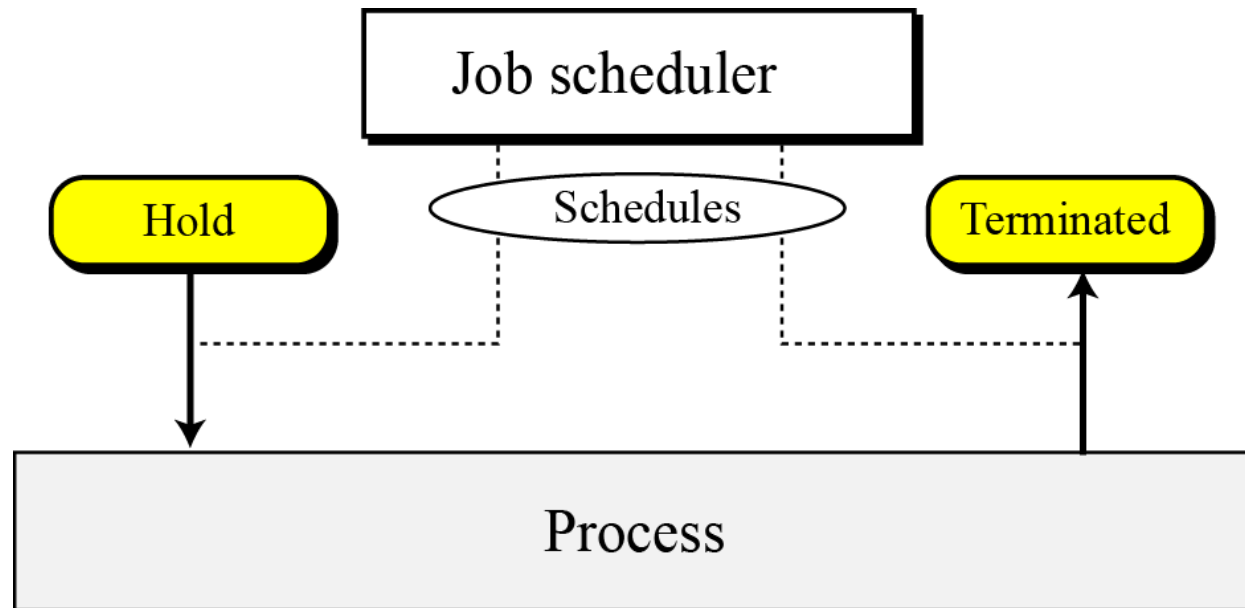
► The relationship can be illustrated with a *state diagram*

1. A program becomes a job when selected by the operating system and brought to the *hold state*
2. When there is memory space available to load the program totally or partially, the job moves to the *ready state* and becomes a process
3. It remains in memory and in this state until the CPU can execute it, moving to the *running state*
4. When it encounters I/O it enters *wait state* and it goes into the *terminated state* when finishing



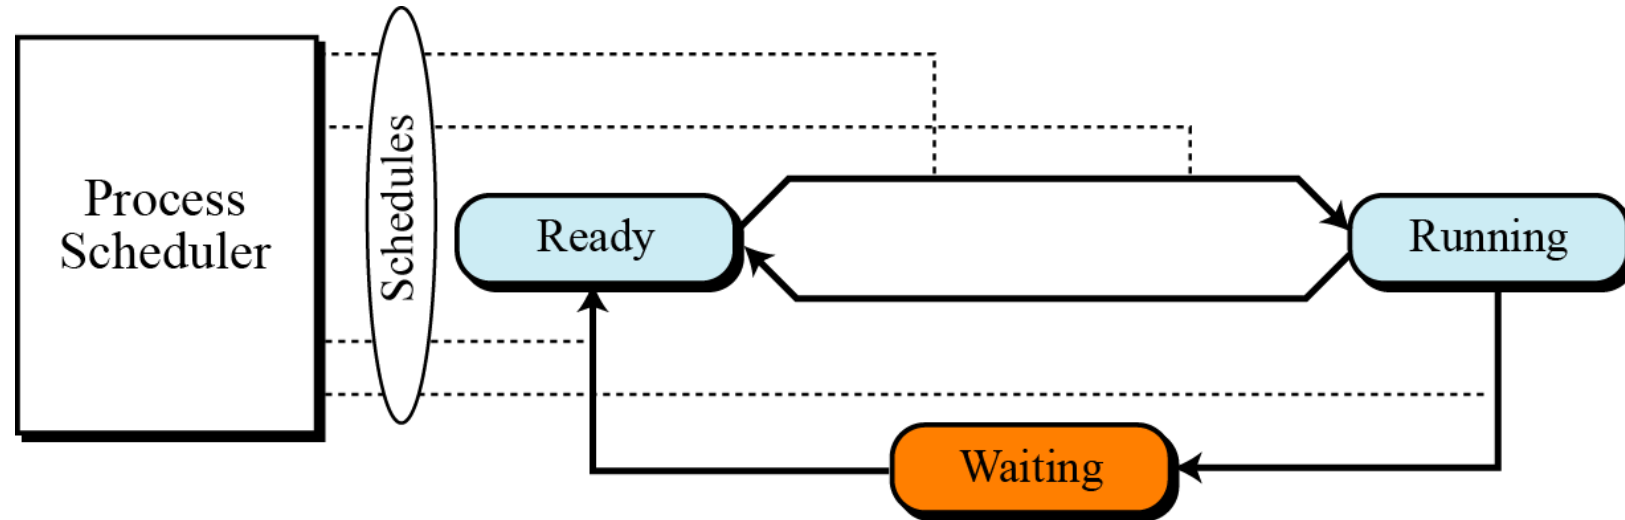
# Process manager - Schedulers

- ▶ To move a job or process from one state to another, the process manager uses two schedulers: the *job scheduler* and the *process scheduler*
  - ▶ A job scheduler is responsible for creating a process from a job and terminating a process



# Process manager - Schedulers

- ▶ The process scheduler moves a process from one state to another
  - ▶ It moves a process from the running state to the waiting state when the process is waiting for some event (interrupt) to happen
  - ▶ It moves a process from the running state to the ready state if the process' time allotment has expired



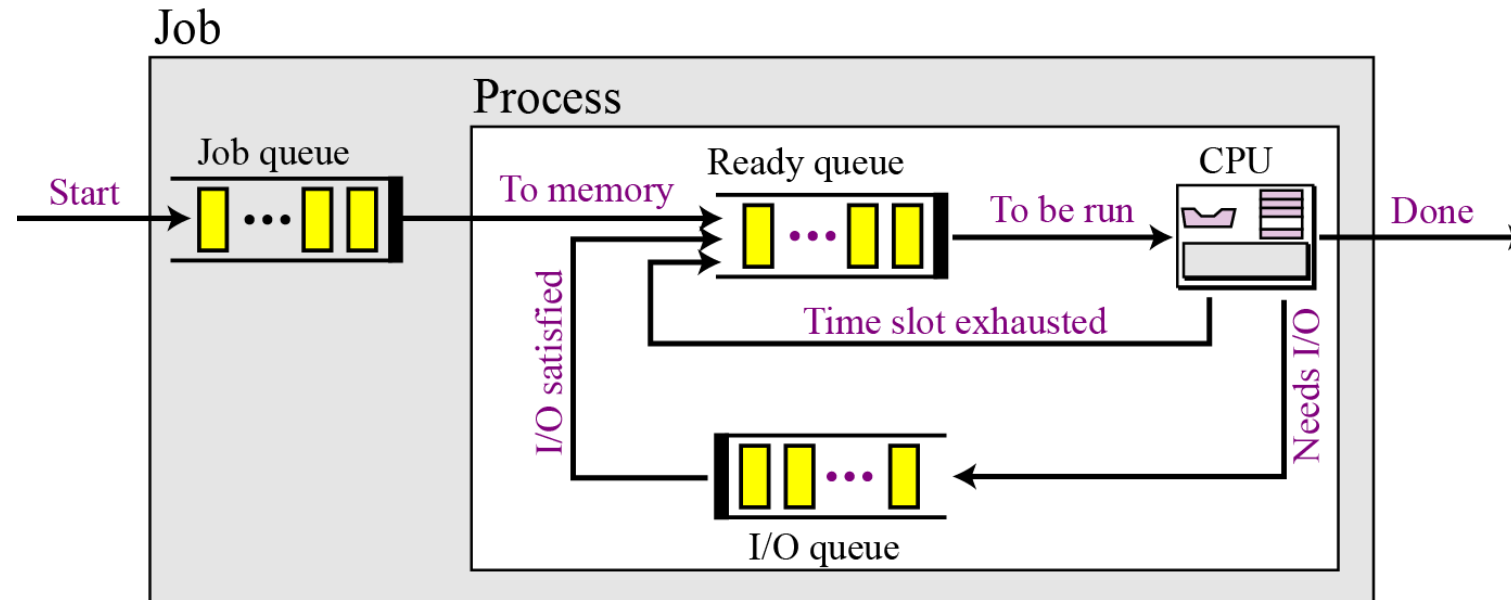
## Process manager - Queuing

---

- ▶ In reality, there are many jobs and many processes competing with each other for computer resources
  - ▶ To handle multiple processes and jobs, the process manager uses *queues*
  - ▶ There is a *job control block* or *process control block* associated with each job or process that stores the information about the job or process
    - ▶ The process manager stores the job or process control block in the queues instead of the job or process itself!


# Process manager - Queuing

- ▶ An OS can have several queue and the process manager can have different policies for selecting the next job or process from a queue



# Process manager - Process synchronization

- ▶ The whole idea behind process management is to synchronize different processes with different resources
  - ▶ We can have two problematic situations: *deadlock* and *starvation*
  - ▶ Deadlock occurs if the operating system allows a process to start running without first checking to see if the required resources are ready, and allows a process to hold a resource as long as it wants



# Process manager - Process synchronization

---

## ▶ Solutions


- ▶ Not to allow a process to start running until the required resources are free
- ▶ Limit the time a process can hold a resource

## ▶ There are four necessary conditions for deadlock

- ▶ Mutual exclusion: Only one process can hold a resource
- ▶ Resource holding: A process holds a resource even though it cannot use it until other resources are available
- ▶ No preemption: The operating system cannot temporarily reallocate a resource (Priority)
- ▶ Circular waiting: All processes and resources involved form a loop


# Process manager - Process synchronization

- ▶ Starvation is the opposite of deadlock. It can happen when the operating system puts too many resource restrictions on a process



a. Process A needs both File1 and File2

b. Process A still needs both File1 and File2



c. Process A still needs both File1 and File2 (starving)



# Device manager

---

- ▶ The device manager, or input/output manager, is responsible for access to input/output devices
  - ▶ The device manager monitors every input/output device constantly to ensure that the device is functioning properly. The manager also needs to know when a device has finished serving one process and is ready to serve the next process in the queue
  - ▶ The device manager maintains a queue for each input/output device or one or more queues for similar input/output devices. For example, if there are two fast printers in the system, the manager can have one queue for each or one queue for both
  - ▶ The device manager controls the different policies for accessing input/output devices. For example, it may use FIFO for one device and the shortest length first for another


# File manager

---

- ▶ File manager is used to controlling the access to files
  - ▶ The file manager controls access to files. Access is permitted only by permitted applications and/or users, and the type of access can vary
  - ▶ The file manager supervises the creation, deletion, modification and naming of files
  - ▶ The file manager supervises the storage of files: how they are stored, where they are stored, and so on
  - ▶ The file manager is responsible for archiving and backups

## Example - UNIX

- ▶ UNIX consists of four major components: the *kernel*, the *shell*, a standard set of *utilities*, and *application programs*
- ▶ Kernel
  - ▶ It contains the most basic parts of the operating system: memory management, process management, device management, and file management
  - ▶ Another component of the kernel consists of a collection of *device drivers*, which are the software units that communicate with the controllers



# Example - UNIX

---

- ▶ Shell

- ▶ It receives and interprets the commands entered by the user

- ▶ Utilities

- ▶ A utility is a standard UNIX program that provides a support process for users
  - ▶ Three common utilities are text editors, search programs, and sort programs

- ▶ Applications

- ▶ Applications in UNIX are programs that are not a standard part of the operating system distribution
  - ▶ They provide extended capabilities to the system



# Appendix

# Memory manager - Demand paging and segmentation


---

- ▶ *Demand paging and segmentation* can be combined to further improve the efficiency of the system
  - ▶ A segment may be too large to fit any available free space in memory. Memory can be divided into frames, and a module can be divided into pages. The pages of a module can then be loaded into memory one by one and executed

# Evolution - Batch systems

---

- ▶ Batch operating systems were designed in the 1950s to control mainframe
  - ▶ At that time, computers used punched cards for input, line printers for output, and tape drives for secondary storage media
  - ▶ Each program to be executed was called a *job* and the punched cards were fed into the computer by an operator
  - ▶ Operating systems during this era were very simple: they only ensured that all of the computer's resources were transferred from one job to the next!



# Time-sharing systems

---

- ▶ The idea of time sharing arises: resources could be shared between different jobs, with each job being allocated a portion of time to use a resource
  - ▶ One means of implementing time-sharing is to apply the technique called *multiprogramming*
  - ▶ For example, when one program is using an input/output device, the CPU is free and can be used by another program
- ▶ Time-sharing is hidden from the user - each user has the impression that the whole system is serving them exclusively!
  - ▶ The operating system now had to do *scheduling*: allocating resources to different programs and deciding which program should use which resource, and when
  - ▶ The user could directly interact with the system without going through an operator!



# Personal systems

---

- ▶ When personal computers were introduced, there was a need for an operating system for this new type of computer
  - ▶ During this era, single-user operating systems such as DOS (Disk Operating System) were introduced

# Parallel systems

---

- ▶ The need for more speed and efficiency led to the design of parallel systems: multiple CPUs on the same machine
  - ▶ Each CPU can be used to serve one program or a part of a program, which means that many tasks can be accomplished in parallel instead of serially
  - ▶ The operating systems required for this are more complex than those that support single CPUs

# Distributed systems

---

- ▶ A job that was previously done on one computer can now be shared between computers that may be thousands of miles apart by network
  - ▶ A program may need files located in different parts of the world
  - ▶ Operating systems combine features of the previous generation with new duties such as controlling security

# Real-time systems

---

- ▶ A real-time system is expected to do a task within specific time constraints
  - ▶ Examples can be found in traffic control, patient monitoring, or military control systems
  - ▶ The application program can sometimes be an embedded system such as a component of a large system, such as the control system in an automobile
  - ▶ The requirements for a real-time operating system are often different than those for a general-purpose system!

## Example - UNIX

---

- ▶ UNIX is a very powerful operating system with three outstanding features
  1. UNIX is a portable operating system that can be moved from one platform to another without many changes. The reason is that it is written mostly in the C language (instead of a machine language specific to a particular computer system)
  2. UNIX has a powerful set of utilities (commands) that can be combined (in an executable file called a script) to solve many problems that require programming in other operating systems
  3. It is device-independent, because it includes device drivers in the operating system itself, which means that it can be easily configured to run any device

## Example - Linux

---

- ▶ The initial kernel of Linux was similar to a small subset of UNIX and now it mainly contain the following components
  - ▶ Kernel
  - ▶ System libraries
    - ▶ The system libraries hold a set of functions used by the application programs, including the shell, to interact with the kernel
  - ▶ System utilities
    - ▶ The system utilities are individual programs that use the services provided by the system libraries to perform management tasks
- ▶ Linux supports the standard Internet protocols
- ▶ Linux' provides the security aspects defined traditionally for UNIX, such as authentication and access control

# Example - Windows

---

- ▶ The design goals released by Microsoft are
  - ▶ **Extensibility**
    - ▶ Windows is designed as a modular architecture with several layers. The purpose is to let the higher layers to be changed with time without affecting the lower layers
  - ▶ **Portability**
    - ▶ Windows, like UNIX, is mostly written in C or C++ and the code is independent of the machine language of the computer on which it is running
  - ▶ **Reliability**
    - ▶ Windows was designed to handle error conditions including protection from malicious software
  - ▶ **Compatibility**
    - ▶ Windows was designed to run programs written for earlier versions of Windows
  - ▶ **Performance**
    - ▶ Windows was designed to have a fast response time to applications that run on top of the operating system

# Example - Windows

---

- ▶ Windows uses a layered architecture
  - ▶ Hardware abstraction layer (HAL)
    - ▶ Hides hardware differences from the upper layers
  - ▶ Kernel
    - ▶ The kernel is the heart of the operating system. It is an object-oriented piece of software that sees any entity as an object
  - ▶ Executive (Runs in kernel (privileged) mode)
    - ▶ The Windows executive provides services for the whole operating system. It is made up of six subsystems: object manager, security reference monitor, process manager, virtual memory manager, local procedure call facility, and the I/O manager
  - ▶ Environmental subsystems (Runs in the user mode)
    - ▶ These are subsystems designed to allow Windows to run application programs designed for Windows, for other operating systems, or for earlier versions of Windows



# Example - Windows

